

XGBoost and Shapley Values for Olympics

Olexiy Pukhov

4/5/2022

The following code makes all plots and images higher resolution.

```
knitr::opts_chunk$set(dpi = 300)
```

This is an exercise to help me get the skills for my research. I tried using XGBoost on the Olympics data. First, we need to install the many packages we need for this. We can use `install.packages` and then `library`, but using `pacman` needs less code to do this and is faster.

```
if (!require("pacman")) install.packages("pacman")

## Loading required package: pacman

pacman::p_load(pacman, rio, tidyverse, xgboost, caTools, ggplot2,
               Ckmeans.1d.dp,
               DALEXtra, mlr, caret, DiagrammeR, SHAPforxgboost, rmarkdown,
               skimr)
```

Importing Data

Now, let's import the data and then remove the values that are highly correlated with other values. Then, let's look at the structure of the data. Then, let's look at the top 20 rows. There are too many rows, so they spill over to the next page. The `##` represents the row #

```
data <- import("olympicmedals.dta")
data = data %>%
  select(-c(cc, year, lpop, lrgdpepc, sptinc992j_p90p100,
            sptinc992j_p99p100))
skimr::skim(data)
```

Data summary

Name	data
Number of rows	204
Number of columns	9

Column type frequency:

numeric	9
---------	---

Group variables	None
-----------------	------

Variable type: numeric

skim_var iable	n_mi ssing	comple te_rate	mean	sd	p0	p25	p50	p75	p100	hist
sptinc99 2j_p0p50	33	0.84	0.15	0.04	0.03	0.12	0.15	0.18	0.26	— ■ ■ ■
rgdpe	28	0.86	7123 59.86	23689 85.23	784. 53	272 90.0 8	9700 8.34	4608 00.94	208605 06.00	■ — — —
pop	28	0.86	43.07	153.75	0.03	2.81	9.72	30.77	1433.7 8	■ — — —
avh	138	0.32	1849. 98	269.24	138 0.61	165 0.92	1818. 28	2061. 05	2474.9 1	■ ■ ■ ■
hc	61	0.70	2.71	0.71	1.22	2.13	2.77	3.26	4.35	■ ■ ■ ■
rnna	29	0.86	3209 645.1 4	10320 354.02	283 3.90	911 49.5 9	3564 59.34	1935 716.6 2	101703 024.00	■ — — —
LifeLadder	55	0.73	5.51	1.10	2.38	4.74	5.49	6.23	7.89	— ■ ■ ■
points	0	1.00	10.28	28.75	0.00	0.00	0.00	8.00	232.00	■ — — —
rgdpepc	28	0.86	2221 3.10	22095. 72	251. 32	510 6.27	1393 3.31	3325 2.37	112941 .45	■ — — —

head(data,n = 20)

```
##      sptinc992j_p0p50      rgdpe      pop      avh      hc      rnna
## 1              NA      3921.261    0.106314      NA      NA      18427.50
## 2              0.1776              NA              NA      NA      NA              NA
## 3              0.0904      228151.016    31.825295      NA      1.481984      1367457.88
## 4              0.1886      35890.020    2.880917      NA      2.964992      227804.64
## 5              NA              NA              NA      NA      NA              NA
## 6              0.1274      681525.812    9.770529      NA      2.746695      4506529.00
## 7              0.1623      991646.312    44.780677      1609.069      3.096804      3399148.50
## 8              0.1854      41048.629    2.957731      NA      3.135995      98812.77
## 9              NA              NA              NA      NA      NA              NA
## 10             NA      1986.163    0.097118      NA      NA      10993.36
## 11             0.1616      1280843.250    25.203198      1726.798      3.549666      5913514.00
## 12             0.2202      498022.250    8.955102      1611.374      3.381046      2878110.75
## 13             0.2029      162126.188    10.047718      NA      NA      275832.31
## 14             0.1402      8664.988    11.530580      NA      1.416526      18935.97
## 15             0.2040      589449.125    11.539328      1586.431      3.149034      3498440.75
## 16             0.1142      36740.742    11.801151      NA      1.918610      88271.80
## 17             0.1504      43608.875    20.321378      NA      1.286242      86734.51
## 18             0.1706      756355.562    163.046161      2418.883      2.101790      2844179.25
```

```
## 19      0.1650 159419.969 7.000119 1645.246 3.186015 456857.28
## 20      0.1014 74097.227 1.641172      NA 2.229507 429310.66
##      LifeLadder points      rgdpepc
## 1      NA      0 36883.7695
## 2      2.375092      0      NA
## 3      NA      0 7168.8579
## 4      5.364910      0 12457.8457
## 5      NA      0      NA
## 6      6.458392      0 69753.2188
## 7      5.900567      4 22144.5137
## 8      5.488087      6 13878.4189
## 9      NA      0      NA
## 10     NA      0 20451.0312
## 11      7.137368     87 50820.6641
## 12      7.213489     10 55613.2422
## 13      5.173389     10 16135.6230
## 14      3.775283      0  751.4789
## 15      6.838761     14 51081.7539
## 16      4.407746      0  3113.3186
## 17      4.740893      1  2145.9604
## 18      5.279987      0  4638.9043
## 19      5.597723     13 22773.8945
## 20      6.173176      2 45148.9727
```

Let's set the same seed for random number generation for reproducibility. Then, let's split the data into a training and testing set. The model will be tested on the training set, and then tested on the testing set for accuracy. 75% of the data is going into the training set, and 25% of the data is going into the testing set.

```
set.seed(1)
split = sample.split(data$points, SplitRatio = 0.75)
training_set = subset(data, split == TRUE)
testing_set = subset(data, split == FALSE)
```

The XGBoost algorithm requires that the input values are in a matrix. The algorithm only accepts numerical data. This also means that if there is categorical variables, we have to change them to dummy variables where each column is either a 0 or 1 depending on the category and there are n - 1 columns to represent the categorical data. In this dataset, there are no categorical values.

Let's change the training and testing sets into a matrix.

```
training_set = training_set %>%
  as.matrix()

testing_set = testing_set %>%
  as.matrix()
```

Model Generation

Let's make the XGBoost model, predict some new values and calculate the error in terms of MSE, MAE and RMSE on the testing set.

```
model = xgboost(data = training_set[,-8], label = training_set[,8], nrounds = 40)

## [1] train-rmse:26.374670
## [2] train-rmse:20.696814
## [3] train-rmse:16.510298
## [4] train-rmse:13.157394
## [5] train-rmse:10.600866
## [6] train-rmse:8.590517
## [7] train-rmse:7.086412
## [8] train-rmse:5.897498
## [9] train-rmse:4.995760
## [10] train-rmse:4.268226
## [11] train-rmse:3.628120
## [12] train-rmse:3.085234
## [13] train-rmse:2.665470
## [14] train-rmse:2.329151
## [15] train-rmse:2.062102
## [16] train-rmse:1.835249
## [17] train-rmse:1.636155
## [18] train-rmse:1.482885
## [19] train-rmse:1.340541
## [20] train-rmse:1.242759
## [21] train-rmse:1.168821
## [22] train-rmse:1.025638
## [23] train-rmse:0.927718
## [24] train-rmse:0.834716
## [25] train-rmse:0.776352
## [26] train-rmse:0.731941
## [27] train-rmse:0.682258
## [28] train-rmse:0.635137
## [29] train-rmse:0.569089
## [30] train-rmse:0.540080
## [31] train-rmse:0.487153
## [32] train-rmse:0.438523
## [33] train-rmse:0.405070
## [34] train-rmse:0.364107
## [35] train-rmse:0.315020
## [36] train-rmse:0.277065
## [37] train-rmse:0.258927
## [38] train-rmse:0.232593
## [39] train-rmse:0.215684
## [40] train-rmse:0.200511

pred_y = predict(model,testing_set[,-8])
```

```

mse = mean((testing_set[8] - pred_y)^2)
mae = caret::MAE(testing_set[8], pred_y)
rmse = caret::RMSE(testing_set[8], pred_y)

cat("MSE: ", mse, "MAE: ", mae, " RMSE: ", rmse)

## MSE: 32.07732 MAE: 3.330121 RMSE: 5.663685

```

Let's compare this to the initial linear model that I made before.

```

data2 <- import("olympicmedals.dta")
data2 <- data2 %>%
  filter(!is.na(data2$lpop))

data2 <- data2 %>%
  filter(!is.na(data2$lrgdpepc))

set.seed(1)
split2 = sample.split(data2$points, SplitRatio = 0.75)
training_set2 = subset(data2, split2 == TRUE)
testing_set2 = subset(data2, split2 == FALSE)

model2 <- lm(points ~ lpop + lrgdpepc, data=training_set2)
pred_y2 = predict(model2, testing_set2[, -12])

mse2 = mean((testing_set2[, 12] - pred_y2)^2)
mae2 = caret::MAE(testing_set2[, 12], pred_y2)
rmse2 = caret::RMSE(testing_set2[, 12], pred_y2)

cat("MSE: ", mse2, "MAE: ", mae2, " RMSE: ", rmse2)

## MSE: 337.4018 MAE: 15.62729 RMSE: 18.3685

```

The XGBoost model is much more accurate. It has an RMSE of 5.66, compared to the linear regression model which has a RMSE of 18.3685.

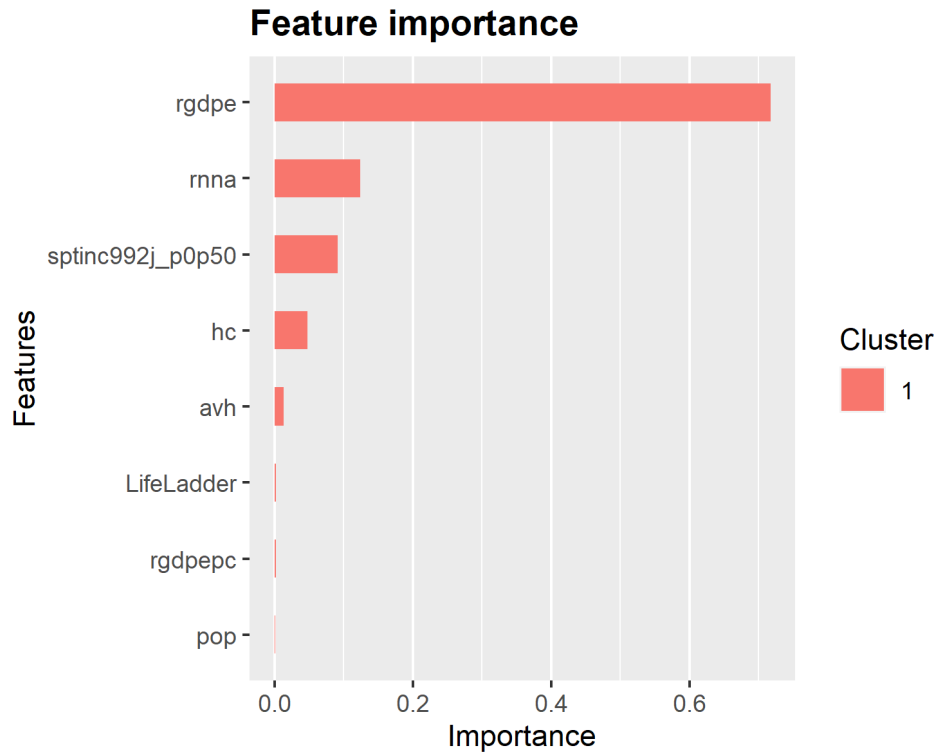
Visualizing the results

Let's now visualize what features are important in this model. Out of all the variables we supplied to the model, which were the most important?

```

xgb_imp <- xgb.importance(feature_names = model$feature_names,
                          model = model)
xgb.ggplot.importance(xgb_imp, n_clusters = 1)

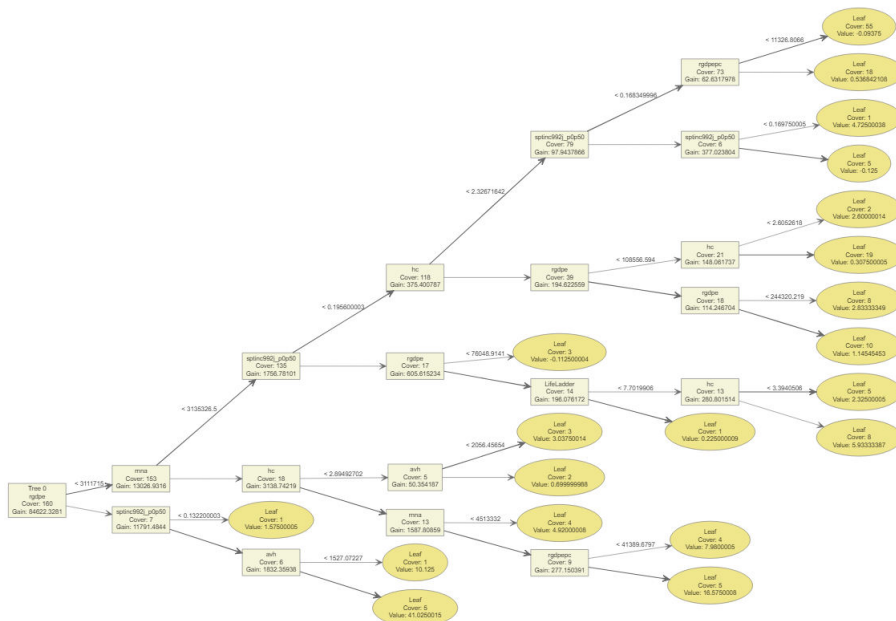
```



It looks like rgdpe (Expenditure-side real GDP at chained PPPs (in mil. 2017 US\$) was the most important variable in predicting the results.

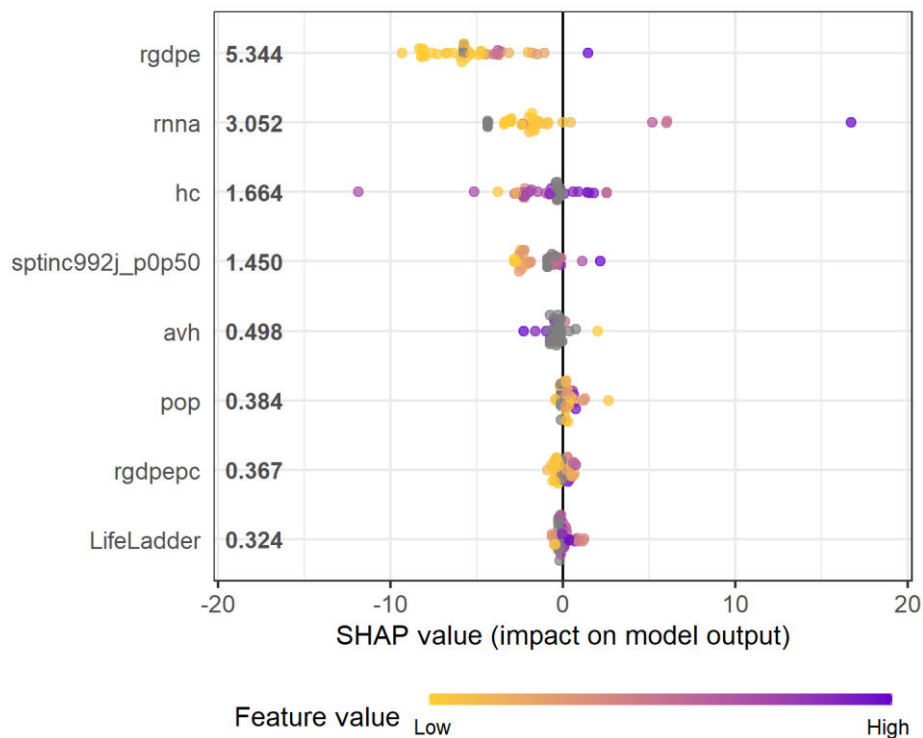
This will plot the first decision tree, which is not really useful for interpretation as XGBoost is an ensemble decision tree model, composed of many trees. However, it is still useful to give us an idea of what the algorithm is thinking about.

```
xgb.plot.tree(model = model, trees = 0)
```



Now let's calculate the SHAP values and see which variable features were most important for our model.

```
shap <- shap.prep(xgb_model = model, X_train = testing_set[, -8])
shap.plot.summary(shap)
```



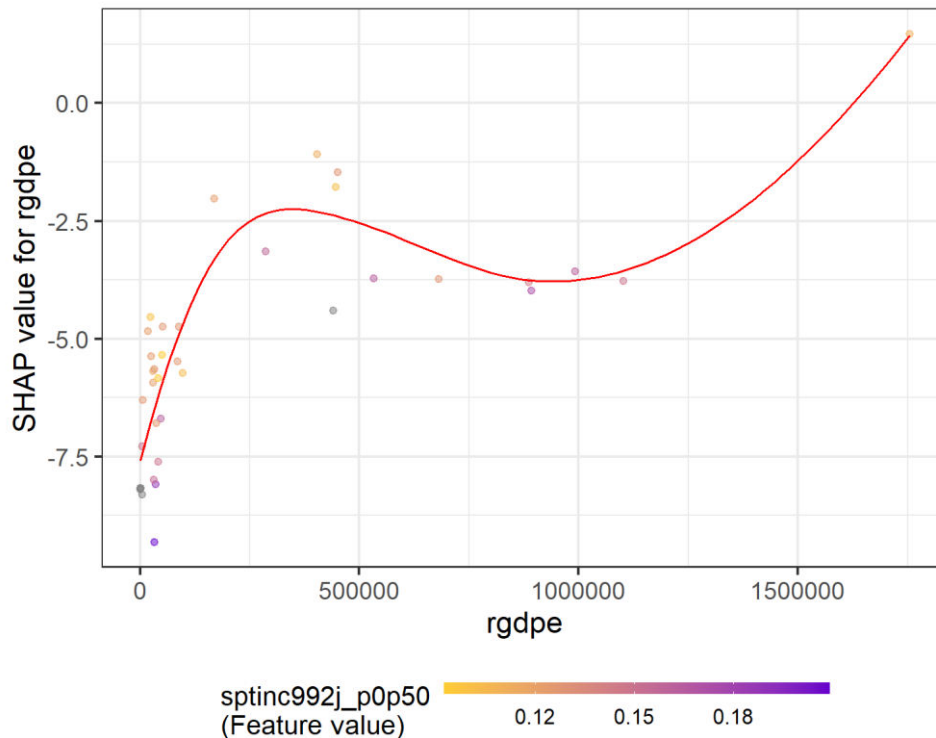
Rgdpe is the most important variable for predicting points won in the olympics. Let's look at this variable in a partial dependence plot.

```
shap.plot.dependence(shap, "rgdpe", color_feature = "auto",
                      alpha = 0.5, jitter_width = 0.1)
```

```
## `geom_smooth()` using formula 'y ~ x'
```

```
## Warning: Removed 7 rows containing non-finite values (stat_smooth).
```

```
## Warning: Removed 7 rows containing missing values (geom_point).
```



It seems that a low `rgdpe` (Expenditure-side real GDP at chained PPPs (in mil. 2017 US\$) hampers your ability to get points at a low value, has less an effect at higher values, and at very high values allows you to obtain higher points. Let us make sure of the results by looking again at our testing set.

```
max(testing_set[,8])
## [1] 10
```

It seems the country with the maximum amount of points in the testing set was 10. If we compare this to our training set:

```
max(training_set[,8])
## [1] 232
```

The maximum amount of points of a country in the training set is 232. This is an result of the countries randomly being put in the training and testing set. How do the results differ if a country with higher points was randomly chosen to be in the testing set? Let's try a `SplitRatio` of 0.7.

```
data <- import("olympicmedals.dta")

data = data %>%
  select(-c(cc, year, lpop, lrgdpepc, sptinc992j_p90p100,
sptinc992j_p99p100))

set.seed(1)
split = sample.split(data$points, SplitRatio = 0.7)
```



```

training_set = subset(data, split == TRUE)
testing_set = subset(data, split == FALSE)

training_set = training_set %>%
  as.matrix()

testing_set = testing_set %>%
  as.matrix()

```

Let's check the testing set to see if it has a higher country with max value for points.

```

max(testing_set[,8])

## [1] 68

```

Let's also check the max value for points in the training set to see if it includes a high point scoring country.

```

max(training_set[,8])

## [1] 232

```

Now let's calculate the model again.

```

model = xgboost(data = training_set[, -8], label = training_set[, 8], nrounds =
40)

## [1] train-rmse:27.064226
## [2] train-rmse:21.260931
## [3] train-rmse:17.006104
## [4] train-rmse:13.642632
## [5] train-rmse:10.957189
## [6] train-rmse:8.904371
## [7] train-rmse:7.340249
## [8] train-rmse:6.071575
## [9] train-rmse:5.093059
## [10] train-rmse:4.339937
## [11] train-rmse:3.710795
## [12] train-rmse:3.165035
## [13] train-rmse:2.737292
## [14] train-rmse:2.358058
## [15] train-rmse:2.060940
## [16] train-rmse:1.819282
## [17] train-rmse:1.568922
## [18] train-rmse:1.391671
## [19] train-rmse:1.254977
## [20] train-rmse:1.144986
## [21] train-rmse:0.995366
## [22] train-rmse:0.869612
## [23] train-rmse:0.750286
## [24] train-rmse:0.683142
## [25] train-rmse:0.635431
## [26] train-rmse:0.596431

```

```

## [27] train-rmse:0.561876
## [28] train-rmse:0.494441
## [29] train-rmse:0.451951
## [30] train-rmse:0.425904
## [31] train-rmse:0.397457
## [32] train-rmse:0.375684
## [33] train-rmse:0.330268
## [34] train-rmse:0.319064
## [35] train-rmse:0.286340
## [36] train-rmse:0.274272
## [37] train-rmse:0.263440
## [38] train-rmse:0.257379
## [39] train-rmse:0.232242
## [40] train-rmse:0.210472

pred_y = predict(model,testing_set[,-8])

mse = mean((testing_set[8] - pred_y)^2)
mae = caret::MAE(testing_set[8], pred_y)
rmse = caret::RMSE(testing_set[8], pred_y)

cat("MSE: ", mse, "MAE: ", mae, " RMSE: ", rmse)

## MSE: 374.9142 MAE: 7.262226 RMSE: 19.3627

```

The RMSE of our model is now 19.36, larger than the RMSE for our linear model at 18.37. At smaller values of points, the XGBoost model works better than the linear model. However, at larger values of points, the XGboost model becomes worse than the linear model. Only a few countries won a lot of points in the olympics (USA, Russia, Great Britain, Japan and China). This is probably because of insufficient data - XGBoost is designed for very large data, and does not work well with a few outliers in a very small dataset.